

# Mapping BPMN Diagrams to JIAC Agent Beans

Tobias Küster      Petrus Setiawan Tan      Axel Heßler

## Abstract

Different mappings from processes to agents have been proposed, but using restricted process models or targeting only single agents, none of them is very convincing. This is a pity, since business processes have many notions in common with agents and would be well suited for modelling complex multi-agent systems. In this paper, we combine two of the existing approaches, namely WADE and the mapping from BPMN to JADL, to a mapping from BPMN to JIAC Agent Beans. The resulting Agent Beans are well-structured and extensible, and at the same time account for nearly all of the expressiveness of BPMN.

## 1 Introduction

In recent times, different approaches have been introduced, using business process diagrams or related notations for modelling agents and multi-agent systems [4, 12]. However, none of these approaches is really compelling. Often, very simple workflow models are used, or if a more expressive process modelling notation, such as BPMN [13] is used, only a limited subset of the language is covered. Further, usually only single agents are targeted, while interactions between agents – which could very well be modelled using e.g. BPMN – are not regarded.

This is a pity, since process diagrams – particularly complex notations such as the Business Process Modeling Notation (BPMN) – share many concepts and abstractions with multi-agent systems. Besides the actual workflow, those notations can be used for modelling several participants in a process, as well as their interactions and communication, or their reactions to external events, centring much more on the *what* and less on the *how*. Thus, despite the shortcomings of existing approaches, BPMN and related notations appear to be very well suited for modelling agents and particularly multi-agent systems.

In this paper we take a look at some of the existing approaches – the WADE extension to the JADE agent framework, and the mapping from BPMN to JIAC’s scripting language JADL – and try to combine the strong sides of both into a new approach. The result is a mapping from BPMN to JIAC Agent Beans, following the basic principles of WADE, but using a much more expressive process notation.

This way, JIAC Agent Beans, being the core components of the agents, can be generated easily from BPMN processes. The resulting Java classes are similarly structured and extensible as the JADE classes of WADE, but reflect the whole BPMN process, including communication between agents and event-handling, both as part of the workflow and for triggering the process.

The remainder of this paper is structured as follows: In Section 2 we will give a short introduction to both, the BPMN process modelling language, and the JIAC multi-agent framework. Then, in Section 3 we will have a look at related work, most notably the WADE framework and the mapping from BPMN to JADL. Thereafter, we will describe in detail how BPMN processes can be mapped to semantically equivalent JIAC Agent Beans (Section 4), and how the transformation has been implemented (Section 5). In Section 6, the mapping is illustrated using an example, before we finally discuss our results in Section 7.

## 2 Background

In the following, we will introduce the reader to the BPMN language and the JIAC agent framework, being the domain an co-domain of the mapping proposed in this paper.

### 2.1 BPMN

The Business Process Modeling Notation [13] is a workflow notation which can be used both as a description language for real-world processes, and as a high-level modelling language for computer programs – most prominently BPEL processes. It can be seen as a combination of UML’s Activity Diagrams and Sequence Diagrams, depicting both the actors’ internal processes and their interactions. An example diagram is shown in Figure 1 in Section 6.

BPMN diagrams can be understood at three levels of abstraction:

1. The diagrams are made up of a few easily recognisable elements, i.e. Events (circles), Activities (boxes) and Gateways (diamonds), connected by Sequence- and Message Flows and situated in one or more Pools.
2. These basic elements are further distinguished using sets of marker icons, e.g. Message, Timer, and Error Events, or parallel and exclusive Gateways.
3. Each element features a number of additional attributes, which are hidden from the diagram, but contain all the information that is necessary for automated code generation, e.g. properties and assignments.

Consequently, the essence of a BPMN diagram is easily understood by all business partners, including those who have great knowledge in their domain but little understanding of programming and multi-agent systems. At the same time, BPMN diagrams provide enough information for the generation of executable programs.

BPMN diagrams have a variety of notational elements, making them well suited for the design of distributed systems in general and multi-agent systems in particular. The process diagrams are subdivided into Pools, each representing one Participant in the process. Using Message Flows for communication between Pools, even complex interaction protocols can be modelled clearly. Further, the notation supports features such as event- and error handling, compensation, transactions and *ad-hoc*-behaviour.

While the semantics of some elements of BPMN – particularly those not covered in the official mapping from BPMN to BPEL [13, Appendix A] – are not

very clearly defined, there is an increasing number of approaches describing the semantics of BPMN using e.g. Petri nets [7], and version 2.0 of the specification makes things clearer, too.

The reason why Petri nets are not used in the first place is that while Petri nets have very clear semantics, and basically everything can be expressed as a Petri net, some high-level constructs that are directly supported by BPMN would result in huge, incomprehensible Petri nets.

BPMN is neither the first process modelling notation, nor will it be the last. However, given its high level of adoption in practical process modelling [15] and its relatively detailed execution semantics, it has proven to be a good choice for modelling distributed computing systems.

## 2.2 JIAC

JIAC V (Java Intelligent Agent Componentware, Version 5) is a Java-based multi-agent development framework and runtime environment [9]. Among others, JIAC features communication, tuple-space based memory, transparent distribution of agents and services, and provides support for dynamic reconfiguration in distributed environments, such as component exchange at runtime. Individual JIAC agents are situated within Agent Nodes, i.e. runtime containers, which also provide support for strong migration. The agents' behaviours and capabilities are defined in a number of so-called *Agent Beans*, which are controlled by the agent's life cycle.

One Agent Bean each JIAC agent is equipped with is the *Communication Bean*, allowing agents to send and receive messages to and from other agents or groups of agents (broadcasting to message channels). The messages are not restricted to FIPA messages but can have any serialisable data as payload. Other commonly used Agent Beans are the *Rule Engine Bean*, integrating a Drools<sup>1</sup> rule engine into the agent's memory, and the *Interpreter Bean*, providing an interpreter for the service-oriented scripting language JADL++ [8].

Besides these and other predefined Agent Beans, the programmer is free to add more Beans to the agent. Each Agent Bean can

- implement a number of *life-cycle* methods, which are executed when the agent changes its life-cycle state, such as initialized, or started,
- implement an *execute*-method, which is called automatically at regular intervals once the agent is running,
- attach *observers* to the agent's memory, being called e.g. each time the agent receives a message or a percept is updated, and
- contribute any number of *action*-methods, which are exposed to the directory and can be invoked by other agents or other beans of the same agent.

Using these four mechanisms, it is possible to define all of the agents' capabilities and behaviours. For details on programming JIAC Agent Beans, please refer to the JIAC Programmers' Manual [10].

---

<sup>1</sup>JBoss Drools: <http://www.jboss.org/drools/>

### 3 Related Work

In the following we will discuss several works which are highly relevant to the approach described in this paper: The original mapping from BPMN to BPEL, a mapping from BPMN to JIAC's scripting language JADL, the WADE framework, mapping workflows to JADE behaviours, and GO-BPMN, a combination of BPMN and goal hierarchies.

#### 3.1 Mapping BPMN to BPEL

One of the motivations for developing BPMN was to provide a standardised graphical notation for *BPEL*, the Business Process Executable Language. Consequently, a mapping from BPMN to BPEL is part of the BPMN specification [13, Appendix A], and a number of alternative or extended mappings have been proposed by various other authors (see for example [14]).

In many aspects, the mapping is very straightforward, as it is evident that BPMN was created with the mapping to BPEL in mind: Each Pool is mapped to a BPEL process (which can be deployed as a Web service), and the several events and activities within are mapped to the workflow of the process. The process is made up mostly of Web service calls, assignments and flow control, but can also contain e.g. event handling based on timing and incoming messages. Given a sufficiently detailed BPMN diagram, the resulting BPEL process can be readily executable.

Still, there are enough elements in BPMN for which no mapping to BPEL is given, i.e. BPMN is not just a visualization for BPEL but an individual language – and in fact more expressive than BPEL itself. Amongst the elements which are not mapped to BPEL are somewhat esoteric elements such as the *ad-hoc* subprocess, or the complex gateway, but also many kinds of events and tasks.

#### 3.2 Mapping BPMN to JADL

In prior work of mapping BPMN to JIAC agents, JIAC's service-oriented scripting language *JADL* [8] was selected as the target of the transformation.

Being conceptually close to BPEL, the mapping is similar, and the process can be mapped very directly to different language elements of JADL – for instance, like BPEL, JADL has dedicated language elements for complex actions such as invoking another service, or for sending and receiving messages, making the generated code compact and easy to comprehend. Similar to the mapping to BPEL, each Pool in the BPMN process is mapped to a JADL service, and the service's input parameters and result types are derived from the Pool's start- and end events [12].

Further, for each start event, a Drools rule is created, starting the respective JADL service on the occurrence of the given event (e.g. an incoming message, or a given time).

Finally, for each Participant in the BPMN process, an XML-based agent configuration file is created, setting up the individual agents, each equipped with an Interpreter Bean and Rule Engine Bean, together with the generated JADL services and Drools rules. Alternatively, the JADL services created from the BPMN processes can be deployed to a running JIAC agent, thus dynamically changing its behaviour.

### 3.3 WADE: Workflows for JADE

A different approach, from which some of the concepts in this work have been drawn, is *WADE (Workflows and Agents Development Environment)*, which is an extension to the JADE multi-agent framework [1]. With WADE, certain aspects of the behaviour of a JADE agent can be modelled using a simple workflow notation [4, 3], based on which Java classes are generated. The workflows basically consist of only two elements: Activities and Transitions.

Using the *Wolf* tool [5], JADE behaviour classes can be generated from those workflow models. In the generated Java classes, there is a clear distinction between the workflow (the order of the activities, together with conditions and guards, where required) and the several activities, each of which is mapped to an individual Java method, which can either refer to existing functionalities or be implemented by the developer. Using this separation, generated workflows can easily be altered or extended.

However, the expressiveness of WADE is restricted by the simplistic workflow model, which allows only the most basic workflows to be modelled. While the transitions can be annotated with guards (conditions), it seems impossible to model parallel execution and synchronisation, let alone more advanced concepts such as event handling or messaging. In fact, each workflow diagram covers only the behaviour of an isolated agent; to our knowledge, interactions between agents can not be modelled.

### 3.4 GO-BPMN

In another approach, *GO-BPMN (Goal-oriented BPMN)*, process models are combined with a goal-hierarchy and executed by agents [6]. The authors praise the high flexibility of the system, and the prospects of parallelisation, but they also write that testing the system is difficult due to possible side-effects of the processes regarding other goals [2].

As the name implies, the individual processes (the “leaves” in the goal hierarchy) are described as BPMN processes. However, only a subset of BPMN is used. Particularly, each diagram shows only a single Pool, and thus, as in the case of WADE, no communication and interaction can be modelled, but just the behaviour of a single agent. While using goals for connecting the individual processes is quite promising, in our opinion process diagrams can more efficiently be used at a higher level of abstraction, e.g. for providing an overview of the system as a whole, instead of for isolated behaviours of individual agents.

## 4 Mapping BPMN to JIAC Agent Beans

When the mapping from BPMN to JADL was already quite advanced, it became apparent that, while being well suited for modelling high-level behaviour or services, traditional JIAC Agent Beans are still advantageous – and often necessary – for defining the better part of the agent’s behaviour. Consequently, work was started on a mapping from BPMN to JIAC Agent Beans [16].

The mapping from BPMN to JIAC Agent Beans is conceptually close to WADE [4]: Each Pool in the BPMN diagram is mapped to one Agent Bean class, with one method for the workflow, and one method for each individual activity

of the process.<sup>2</sup> The *workflow-method* acts as an entry point to executing the process, while the several *activity-methods* are invoked by the workflow method in accordance with the ordering of the activities in the process.

## 4.1 Workflow Method

Basically, the workflow method is made up of calls to the several activity-methods, being arranged into sequences, if-else statements and loops. While this requires the process to be structured properly (see Section 5), the result is very readable and understandable, just like manually written code.

At the same time, BPMN allows for much more expressive workflows to be modelled than the rather minimalistic workflow notation used in WADE. In particular, the following concepts of BPMN are covered by the mapping:

- Parallel execution (BPMN’s AND-Gateway) is mapped to multiple threads being started and joined.
- Subprocesses (composite activities) are mapped to internal classes following the same schema as the main class, with workflow- and activity-methods for the activities embedded into the subprocess.
- Event Handler (Intermediate Events attached to an Activity) are also mapped to threads, running concurrently to the thread executing the Activity itself, and interrupting this thread in case the respective event occurs (e.g. a message or a timer). The workflow is routed accordingly.
- The same pattern is applied to Event-based XOR-Gateways. In this case the main thread will wait until one of the events has been triggered.

An example for a complex workflow-method is given later in Section 6.

## 4.2 Properties and Assignments

BPMN specifies a number of non-visual attributes, such as properties (i.e. variables) and assignments. Properties can be declared in the scope of whole Processes or individual activities (both atomic Tasks and composite Subprocesses). When declared in the scope of a process or subprocess, the property is visible to all elements (transitively) contained therein.

Accordingly, properties are mapped to variables in different scopes in the Agent Bean, reflecting their visibility in the BPMN diagram. Properties of the process are mapped to variables in the scope of the Agent Bean class, properties of a subprocess to variables in the scope of the embedded subprocess class, and properties of an activity to local variables in the scope of the activity-method.

Assignments are always bound to an activity or events, and are included in the respective activity-method. In BPMN, assignments can have an *assign-time* of either ‘before’ or ‘after’, determining whether the assignment has to be applied before or after the actual activity it is bound to is executed.

---

<sup>2</sup>In the following, we will use the term “workflow” for the order in which individual activities are executed in the process, and the term “process” for the whole ensemble of activities and their ordering, events, variables etc.

### 4.3 Activity Methods

The several activity-methods have neither parameters nor a return value and always follow the same schema:

1. *Properties*: First, the properties in the scope of the activity are initialized, if any. For each property, one Java variable is created, using a corresponding data type, being visible only in the scope of this activity.
2. *Start Assignments*: Then, assignments of the activity with assign-time ‘before’ are created, e.g. for assigning values to the input parameters of a service call.
3. *Activity Behaviour*: Now, the code corresponding to the actual activity is inserted, e.g. invoking a service, sending a message, or executing a user-defined code-snippet. If the Activity’s *loop* attribute is set, this part is placed inside of a loop.
4. *End-Assignments*: Finally, assignments with assign-time ‘after’ are created, e.g. for binding the return value of a service call to a variable.

Similar to the mapping to JADL, we can make use of JIAC’s communication infrastructure. Likewise, Message Events and Send and Receive Tasks are mapped to sending and receiving JIAC messages, while Service Tasks are mapped to the invocation of a JIAC action (i.e. a service). Script Tasks allow the developer to attach a custom snippet of Java code to the task. Further, Timer Events can be mapped to a temporary suspension of the execution.

There are more types of Tasks and Events in BPMN, for which no mapping has been devised yet, but these are the most common and important ones.

### 4.4 Start Events and Starter Rules

Finally, the processes’ Start Events have to be mapped to mechanisms for starting the process on the occurrence of the respective events. In the mapping to JADL, a number of Drools rules are created for this purpose. Using Agent Beans, these ‘starter rules’ can be integrated directly into the code, making use of the mechanisms introduced in Section 2.2.

- If the process has a Start Event with unspecified type, or *None* type, then the workflow-method is invoked in the Agent Bean’s `doStart()` method (one of the *life-cycle*-methods), being called when the agent is started.
- For a *Timer* Start Event, the Agent Bean is given an `execute()` method, regularly checking the current time against the time the process was last started, invoking the workflow-method at a given time or interval.
- A *Message* Start Event results in a message observer being attached to the agent’s memory when the Agent Bean is started, which will then invoke the workflow-method every time a matching JIAC message is received.
- Finally, in case of a *Service* Start Event, the workflow-method is marked with the `@Expose` annotation, exposing the workflow-method as a JIAC action to be discovered and invoked by other agents.<sup>3</sup>

---

<sup>3</sup>There is, as such, no Service Start Event in BPMN. We use this term to distinguish Message Start Events, where the message is in fact a service request.

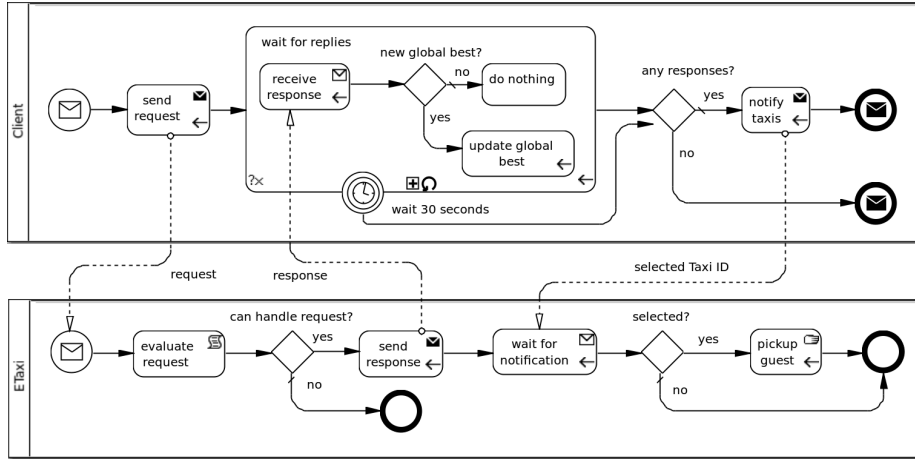


Figure 1: Example: Taxi Request Service

Besides creating these mechanisms, a Service Start Event also results in the workflow-method's input parameters being updated to correspond to the specified service parameters. Analogously, a Service End Event results in the workflow-method's return value being set accordingly.

## 5 Implementation

The mapping has been implemented in the course of the diploma thesis of one of the authors of this paper [16] and is an extension to the BPMN editor *VSDT* (*Visual Service Design Tool*) [11].

The first step in mapping BPMN to Agent Beans – or any structured programming language – is to structure the process graph to a tree of sequences, decision blocks, loops, etc. This functionality is provided by the VSDT's transformation framework, so that only the actual mapping of individual process elements to fragments of Java code, as specified in the last section, had to be implemented.

This element mapping has been separated into two stages. First, the structured process model is translated to an intermediate model, being a high-level representation of the structure of a JIAC Agent Bean. Then, this model can be translated straightforwardly to executable Java code using a number of templates for the JET framework.<sup>4</sup> Using JET and JMerge, parts of the generated code can safely be modified and merged in case the process model changes and has to be re-generated.

## 6 Example

In this section we will exemplify several aspects of the mapping by means of a simple example, as shown in Figure 1.

<sup>4</sup>JET (Java Emitter Templates) is part of the Eclipse Model To Text (M2T) project: <http://www.eclipse.org/modeling/m2t/>



The BPMN diagram consists of two Pools, each representing an agent role: Client, and E-Taxi. The Client's process starts by broadcasting a request (customer ID, location, destination, desired time of arrival) to all available E-Taxis, which evaluate the request and decide whether they can handle it, in which case they send a response (taxi ID, estimated time of arrival, price). The Client enters a looping subprocess, listening to responses and memorizing the best response, until after 30 seconds the subprocess is interrupted by the attached timer event. The Client then sends notifications to the responding E-Taxis, and the selected E-Taxi prepares to pick up the guest. Note that the several properties (variables) and assignments are not visualized in the diagram.

The resulting Agent Bean for the Client role is shown in Listing 1, slightly shortened to improve readability and to better fit into this paper.

Listing 1: Generated Agent Bean: Client\_requestTaxi

```
public class Client_requestTaxi extends AbstractMethodExposingBean {
    public static String ACTION_REQUESTTAXI = "Client_requestTaxi#requestTaxi";
    TaxiResponse globalBest;
    boolean noResponse;
    String appID;
    String location;
    String destination;

    @Expose(name = ACTION_REQUESTTAXI, scope = ActionScope.GLOBAL)
    public String requestTaxi(String currentLoc, String currentDest) {
        location = currentLoc;
        destination = currentDest;
        sendRequest();
        WaitForRepliesSub waitForReplies = new WaitForRepliesSub();
        TimeoutEvent wait30Seconds = new TimeoutEvent(30000, waitForReplies);
        try {
            waitForReplies.start();
            wait30Seconds.start();
            waitForReplies.join();
            wait30Seconds.stop();
        } catch (InterruptedException e) { }
        if (noResponse) {
            String taxiID = "none_available";
            return taxiID;
        } else {
            notifyTaxis();
            String taxiID = globalBest.getTaxiID();
            return taxiID;
        }
    }

    private void sendRequest() {
        TaxiRequest request = new TaxiRequest(location, appID, destination);
        Action sendAction = retrieveAction(ACTION.SEND);
        IGroupAddress group = createGroupAddress("TaxiRequest");
        JiacMessage msg = new JiacMessage(request);
        invoke(sendAction, new Serializable[] {msg, group});
    }

    private void notifyTaxis() {
        [ analogous to sendRequest() ]
    }

    class WaitForRepliesSub extends Thread { [see below] }

    class TimeoutEvent extends Thread { [see below] }
}
```

As can be seen, the control-flow of the process is reflected in the workflow-method `requestTaxi()`, which is also exposed as a JIAC Action, or Service. The workflow-method is dominated by the threads for running the subprocess and the attached event handler, but also contains an if-else-statement for the

Gateway at the end of the process. The activities *send request* and *notify taxis* are mapped to two similar methods for sending a JIAC message to the specified broadcast message group.

The subprocess is mapped to an inner class, also forming a new variable scope for the properties of the subprocess. It features another workflow-method and several activity-methods, most notably the *receive responses* method, in which the Client joins the specified message group and checks its memory for accordant messages. The subprocess is executed as a thread which is to be interrupted by the event handler thread. Both are shown in Listing 2.

Listing 2: Inner Classes for Subprocess and Event Handler

```
class WaitForRepliesSub extends Thread {
    TaxiResponse currentResponse;

    public void run() {
        noResponse = true;
        while (true) {
            receiveResponse();
            if (currentResponse.compareTo(globalBest) > 0) {
                updateGlobalBest();
            } else {
                doNothing();
            }
        }
    }

    private void receiveResponse() {
        Action joinAction = retrieveAction(ACTION_JOIN_GROUP);
        Action leaveAction = retrieveAction(ACTION_LEAVE_GROUP);
        IGroupAddress group = createGroupAddress("TaxiResponseTo"+appID+"");
        invoke(joinAction, new Serializable[] {group});
        TaxiResponse response = null;
        while (response == null) {
            for (JiacMessage msg : memory.readAll(new JiacMessage())) {
                if (msg.getHeader(Header.SEND_TO).equals("TaxiResponseTo"+appID+"")
                    && msg.getPayload() instanceof TaxiResponse) {
                    memory.remove(msg);
                    response = (TaxiResponse) msg.getPayload();
                    break;
                }
            }
            try {
                Thread.sleep(100);
            } catch (InterruptedException e) {}
        }
        invoke(leaveAction, new Serializable[] {group});
        noResponse = false;
        currentResponse = response;
    }

    private void doNothing() {}

    private void updateGlobalBest() {
        globalBest = currentResponse;
    }
}

class TimeoutEvent extends Thread {
    long timeout;
    Thread toStop;
    boolean triggered = false;

    public TimeoutEvent(long timeout, Thread toStop) {
        this.timeout = timeout;
        this.toStop = toStop;
    }

    public void run() {
```

```

    try {
        Thread.sleep(timeout);
        triggered = true;
        toStop.stop();
    } catch (InterruptedException e) { }
}

```

## 7 Conclusion

In this paper, we have presented a new approach of creating multi-agent systems from process models, combining the mapping from BPMN to JADL [12] with ideas borrowed from WADE [4]. The result is a transformation from BPMN process diagrams to JIAC Agent Beans, comprising one method for the workflow as a whole, and one method for each individual activity, but also supporting beneficial aspects of BPMN such as messaging and event handling.

### 7.1 Discussion

Using the domain-specific scripting language JADL, agent behaviours can be expressed in a very compact and readable way, but the overall expressiveness is limited. JIAC Agent Beans, on the other hand, have the full expressiveness of the Java language at their disposal. Thus, basically everything that can be modelled in a BPMN diagram can be mapped to an Agent Bean.

While the resulting workflow-methods for complex processes can become somewhat bulky – particularly if event handlers are used – the separation into workflow-methods and activity-methods keeps the resulting code reasonably clear. Further, like in WADE, individual activity-methods can be altered or extended without risk of losing the changes after the code is generated anew.

One drawback compared to the mapping to JADL scripts is that the generated Agent Beans – being Java classes – can not as easily be deployed to an agent at runtime. Regarding the high expressiveness of the generated Agent Beans and the good performance of compiled Java code compared to the interpreted JADL scripts, the mapping from BPMN to JIAC Agent Beans is suited best for modelling and generating core components of the agents, while the mapping to JADL is of much use for creating dynamic behaviours and services to be deployed and changed at runtime.

### 7.2 Future Work

While the mapping can already be used for generating useful agent behaviours, it is not yet completed. First, there are still aspects of BPMN, which are not covered by the mapping, such as some of the less common event types. Second, there are aspects of agents that can not yet be modelled adequately with BPMN.

Among the latter are messages to individual agents. Originally, BPMN knew only services and service calls. In the course of our work, we extended our ‘dialect’ of BPMN with broadcast message groups, but what’s still missing are messages to individual agents, as identified by their agent ID.

Another issue which we want to tackle in the future is the modelling of goals by means of BPMN. One approach that appears to be promising is to use the *ad-hoc* subprocess for this task, but this is still work in progress.

## References

- [1] F. Bellifemine, A. Poggi, and G. Rimassa. JADE – a FIPA-compliant agent framework. Internal technical report, Telecom Italia, 1999. Part of this report has been also published in Proceedings of PAAM’99, London, April 1999, pp.97-108.
- [2] B. Burmeister, M. Arnold, F. Copaciu, and G. Rimassa. BDI-agents for agile goal-oriented business processes. In *Proc. of 7th Int. Conf. on Autonomous Agents and Multiagent Systems (AAMAS 2008)*, pages 37–44, Richland, SC, 2008. International Foundation for Autonomous Agents and Multiagent Systems.
- [3] G. Caire. *WADE User Guide, Version 2.6*. Telecom Italia, July 2010. <http://jade.tilab.com/wade/doc/WADE-User-Guide.pdf>.
- [4] G. Caire, D. Gotta, and M. Banzi. WADE: A software platform to develop mission critical applications exploiting agents and workflows. In M. Berger, B. Burg, and S. Nishiyama, editors, *Proc. of 7th Int. Conf. on Autonomous Agents and Multiagent Systems (AAMAS 2008) – Industry and Applications Track*, pages 29–36, May 2008.
- [5] G. Caire, E. Quarantotto, M. Porta, and G. Sacchi. WOLF – an eclipse plug-in for WADE. In *Proceedings of the ACEC 2008*, 2008.
- [6] M. Calisti and D. Greenwood. Goal-oriented autonomic process modeling and execution for next generation networks. In *3rd IEEE International Workshop on Modelling Autonomic Communications Environments (MACE)*. Springer, 2008.
- [7] R. M. Dijkman, M. Dumas, and C. Ouyang. Formal semantics and analysis of BPMN process models using Petri nets, 2008.
- [8] B. Hirsch, T. Konnerth, M. Burkhardt, and S. Albayrak. Programming service oriented agents. In M. Calisti, F. P. Dignum, R. Kowalczyk, F. Leymann, and R. Unland, editors, *Service-Oriented Architecture and (Multi-)Agent Systems Technology*, number 10021 in Dagstuhl Seminar Proceedings, Dagstuhl, Germany, 2010. Schloss Dagstuhl – Leibniz-Zentrum fuer Informatik, Germany.
- [9] B. Hirsch, T. Konnerth, and A. Heßler. Merging agents and services – the JIAC agent platform. In R. H. Bordini, M. Dastani, J. Dix, and A. El Fallah Seghrouchni, editors, *Multi-Agent Programming: Languages, Tools and Applications*, pages 159–185. Springer, 2009.
- [10] JIAC Development Team. *JIAC – Java Intelligent Agent Componentware, Version 5.1.0*. DAI-Labor, TU Berlin, July 2010. <http://www.jiac.de>.
- [11] T. Küster and A. Heßler. Towards transformations from BPMN to heterogeneous systems. In D. Ardagna, M. Mecella, and J. Yang, editors, *Business Process Management Workshops*, volume 17 of *LNBIP*, Milano, Italy, September 2008. Springer.

- [12] T. Küster, M. Lützenberger, A. Heßler, and B. Hirsch. Integrating process modelling into multi-agent system engineering. In M. Huhns, R. Kowalczyk, Z. Maamar, R. Unland, and B. Vo, editors, *Proceedings of the 5th Workshop of Service-Oriented Computing: Agents, Semantics, and Engineering (SOCASE)*, 2010.
- [13] Object Management Group. Business process modeling notation (BPMN) version 1.2. Specification formal/2009-01-03, Object Management Group, January 2009.
- [14] C. Ouyang, M. Dumas, W. M. P. van der Aalst, A. H. M. ter Hofstede, and J. Mendling. From business process models to process-oriented software systems. *ACM Transactions on Software Engineering and Methodology*, 19(1):1–37, August 2009. New York, NY, USA.
- [15] J. C. Recker. BPMN modeling – who, where, how and why. *BPTrends*, 5(3):1–8, March 2008.
- [16] P. S. Tan. Automated generation of JIAC agentbeans from BPMN diagrams. Diplomarbeit, Technische Universität Berlin, November 2011.